

Using Closed-Loop Detection to Improve Homography Estimation and Mosaicing

Michele Pratusевич

July 25, 2012

Contents

1	Introduction and Overview	1
2	Homographies	2
2.1	In General	2
2.2	Factoring a Homography Matrix	2
2.3	Computing the Homography Matrices in Code: Overview	2
3	Detecting a Closed Loop	3
4	Optimization Problem Formulation	3
5	Results	5
6	Discussion	5
7	Potential Further Work	5
8	Conclusion	5
A	Homography Calculation Code	6
B	Parameters in OpenCV Methods	6

1 Introduction and Overview

The goal of this project is to improve an automatically mosaiced image by detecting a closed loop in a sequence of images and re-computing the calculated image homographies using nonlinear optimization. First homographies will be discussed, with some discussion about factoring homography matrices to determine the various components of a homography. Next the techniques for detecting a closed loop in an image sequence are discussed. Then the optimization problem for improving the homography matrices is described in terms of

formulation and the constraints. Lastly, further considerations and ideas are discussed.

2 Homographies

2.1 In General

A homography H in this context is a 3 by 3 matrix $H_{i,j}$ that describes the relationship between image i and image j in an image sequence. In this application, the homographies $H_{i,i+1}$ are calculated between consecutive images i and $i+1$. When calculating an image mosaic, the cumulative homography for image k in the sequence is calculated by multiplying each of the previous homography matrices together until the current image:

$$H_{1,k} = H_{1,2} \cdot H_{2,3} \cdots H_{k-1,k}$$

In the application of creating a 2-D mosaic from a video, the homography is a 2D homography, transforming a point $[u, v, 1]^T$ to $[u', v', 1]^T$ up to a scale factor, normalized in OpenCV to be 1.

2.2 Factoring a Homography Matrix

As discussed by Sonka, et. al in [1], homographies form a group under multiplication as shown above. There are various important subgroups that can be used to factor any homography matrix H into its components. As discussed in [1], any homography can be decomposed as $H = H_P H_A H_S$ where

$$H_P = \begin{bmatrix} I & \mathbf{0} \\ \mathbf{a}^T & b \end{bmatrix}, \quad H_A = \begin{bmatrix} K & \mathbf{0} \\ \mathbf{0}^T & 1 \end{bmatrix}, \quad H_S = \begin{bmatrix} R & -R\mathbf{t} \\ \mathbf{0}^T & 1 \end{bmatrix}$$

And K is upper triangular. Page 557 in [1] gives more information about these subgroups.

Using the factorization above as a starting point, a homography matrix is therefore written as follows:

$$H = H_P H_A H_S = \begin{bmatrix} KR & -KR\mathbf{t} \\ \mathbf{a}^T KR & -\mathbf{a}^T KR\mathbf{t} + b \end{bmatrix}$$

Where \mathbf{t} is the translation vector, R is the rotation matrix for rotation in the xy -plane, and K is the dilation matrix in the xy -directions. \mathbf{a} and b are coefficients related to the rotation and dilation in the z direction.

2.3 Computing the Homography Matrices in Code: Overview

The homography matrices can be computed using an OpenCV method called `cvFindHomography`, described in [?], which needs to be tuned according to the application you are working with. For instance, the window size that is needed

to find a correct homography depends on how close the camera is to the scene in question.

To use the `cvFindHomography` method you already need to have points picked out on both images that are known to be correspondent. In applications with few images, these points can be picked out by hand and visual inspection. Because the goal of this project was to create a fully automated mosaicing algorithm, manually picking points was not an option. Instead, because the image sequence is taken from a video, the video can be sequenced in such a way to make the images overlapped by a significant amount. Given this information, we use optical flow in the OpenCV method `cvCalcOpticalFlowPyrLK` to calculate the new positions of interest points from the first image in the second image, described in [4]. The optical flow method has parameters to tune as well.

The interest points in the first image are found using the OpenCV methods `cvGoodFeaturesToTrack` and `cvFindCornerSubPix`, described in [3]. The parameters on these methods must be tuned as well.

The `cvFindHomography` method then uses correspondences in those points to compute an image-wide homography that is then used to construct the mosaic.

The code is given in A and the information about adjusting parameters in all the functions is given in B.

3 Detecting a Closed Loop

Based on the factorization of a homography discussed in ??, various components of a homography matrix can be factored out and used in an algorithm to detect a closed loop.

TODO

4 Optimization Problem Formulation

This is the problem formulation for the optimization problem in closed-loop homographies.

There is a closed loop detected between image 1 and image n . Each homography computed by the mosaicing algorithm, $H_{i,i+1}$, is the homography from image i to the next image $i+1$, computed using OpenCV (and RANSAC). Once the closed loop is detected, the homography $H_{1,n}$ is computed between image 1 and the overlap image n .

Overall, the goal of the optimization algorithm is to minimize the error between

$$H_{1,2} \cdot H_{2,3} \cdots H_{n-2,n-1} \cdot H_{n-1,n} - H_{1,n} = 0$$

Which can be written as

$$H_{1,n}^{\text{cumulative}} - H_{1,n} = 0.$$

Using a scalar formulation of the problem, the error is calculated by calculating the sum of the absolute values of the differences between the cumulative and new homographies for each component in the matrix.

$$\sum_{i,j} |(H_{1,n}^{\text{cumulative}})_{ij} - (H_{1,n})_{ij}|.$$

The variables given to the optimizer are the 8 parameters in each of the matrices $H_{i,i+1}$ (each entry of the homography matrix except for the 3,3 entry, which is assumed to be 1. The values of the $H_{1,n}$ matrix are not variables in the optimization function, but are taken as truth.¹ The optimizer will yield n new homography matrices $H_{i,i+1}^{\text{optimized}}$.

The optimizer tries to minimize this nonlinear multivariable function according to the following constraints:

1. The new matrices that are computed must be homography matrices (i.e. still have determinant close to 1).

$$|\det(H_{i,i+1}) - 1| \leq \text{determinant_threshold} \quad \forall i \in (1, n-1)$$

2. The variables in each matrix can't change "too much." There are a few ways to implement this:

- (a) The individual variables can only change within a certain range (with this range being different depending on which component in the homography it is):

$$\begin{aligned} |(H_{i,i+1})_{1,1} - (H_{i,i+1}^{\text{optimized}})_{1,1}| &\leq \text{change_threshold} \\ |(H_{i,i+1})_{1,2} - (H_{i,i+1}^{\text{optimized}})_{1,2}| &\leq \text{change_threshold} \\ |(H_{i,i+1})_{2,1} - (H_{i,i+1}^{\text{optimized}})_{2,1}| &\leq \text{change_threshold} \\ |(H_{i,i+1})_{2,2} - (H_{i,i+1}^{\text{optimized}})_{2,2}| &\leq \text{change_threshold} \\ |(H_{i,i+1})_{3,1} - (H_{i,i+1}^{\text{optimized}})_{3,1}| &\leq \text{pixel_threshold} \\ |(H_{i,i+1})_{3,2} - (H_{i,i+1}^{\text{optimized}})_{3,2}| &\leq \text{pixel_threshold} \\ |(H_{i,i+1})_{1,3} - (H_{i,i+1}^{\text{optimized}})_{1,3}| &\leq \text{small_threshold} \\ |(H_{i,i+1})_{2,3} - (H_{i,i+1}^{\text{optimized}})_{2,3}| &\leq \text{small_threshold} \end{aligned}$$

The reason each entry (or group of entries) should have it's own threshold is that each component of the homography matrix is related to a different transformation and has different similarity tolerances.

¹It is possible to use the components of the closed-loop matrix as variables too, but for now I've implemented the code to consider it as truth. I will try using those components as variables as well.

- (b) The total change in all the variables of a particular matrix can't exceed a certain value.

$$|\sum_{i,j}(H_{k,k+1})_{i,j}| \leq \text{sum_thresh}$$

I don't like this way because it does not account for two very large changes in two variables - i.e. if component (1,1) changes by -100 and component (3,3) changes by $+99$, it seems that there wasn't that much change in total whereas in reality the new homography matrix is very different from the old homography matrix.

- (c) Factor the new homography matrix and allow not "too much" of a change from the new translation and rotation components.
- (d) Don't care about changes between each of the individual homographies but care about the changes in all the cumulative homographies.

Right now my code uses the 2a metric of "too much change."

3. The 3,3 entry of all the intermediate homographies computed with $H^{\text{optimized}}$ cannot be very different from 1.

$$\text{Let } H_{1,k}^{\text{optimized}} = H_{1,2}^{\text{optimized}} \dots H_{k-1,k}^{\text{optimized}}$$

$$|(H_{1,k}^{\text{optimized}})_{3,3} - 1| \leq \text{entry33_threshold } \forall k \in (2, n)$$

5 Results

TODO

6 Discussion

TODO

7 Potential Further Work

masks for detecting features
breaking up the image TODO

8 Conclusion

TODO

References

- [1] M. Sonka, V. Hlavac, and R. Boyle, *Image Processing, Analysis, and Machine Vision*, 3rd ed., Thomson, USA, 2008.
- [2] http://opencv.willowgarage.com/documentation/camera_calibration_and_3d_reconstruction.html
- [3] http://opencv.willowgarage.com/documentation/feature_detection.html
- [4] http://opencv.willowgarage.com/documentation/c/video_motion_analysis_and_object_tracking.html

A Homography Calculation Code

The following is code to compute the features to track in the first image.

```
IplImage* imgfirstBW = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
CvPoint2D32f* cornersA = new CvPoint2D32f[MAX_CORNERS];
IplImage* eig_image = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
IplImage* tmp_image = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
const int MAX_CORNERS = 500;
int corner_count = MAX_CORNERS;

//not shown: saving image to imgfirstBW

//compute the features to track from the first image
cvGoodFeaturesToTrack(imgfirstBW, eig_image, tmp_image, cornersA, &corner_count, 0.01, 5,
cvFindCornerSubPix(
    imgfirstBW,
    cornersA,
    corner_count,
    cvSize(10, 10),
    cvSize(-1, -1),
    cvTermCriteria(CV_TERMCRIT_ITER|CV_TERMCRIT_EPS, 20, 0.03));
```

This code is used to calculate the optical flow between two images and convert the computed array of points into a matrix that can be used for the homography method.

```
IplImage* imgBW = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
IplImage* pyr1 = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
IplImage* pyr2 = cvCreateImage(cvSize(imgWidth, imgHeight), IPL_DEPTH_8U, 1);
CvPoint2D32f* cornersB = new CvPoint2D32f[MAX_CORNERS];
char features_found[MAX_CORNERS];
float feature_errors[MAX_CORNERS];
```

```

CvMat* PointImg1;
CvMat* PointImg2;

//not shown: saving the image to imgBW

//compute homography with first image and the new image
cvCalcOpticalFlowPyrLK(
    imgfirstBW,
    imgBW,
    pyr1,
    pyr2,
    cornersA,
    cornersB,
    corner_count,
    cvSize(30, 30),
    3,
    features_found,
    feature_errors,
    cvTermCriteria(CV_TERMCRIT_ITER | CV_TERMCRIT_EPS, 20, 0.03),
    0
);

countfound = 0;

//count how many matched features you get from the optical flow
for (int i = 0; i < corner_count; i++)
{
    if (features_found[i] == 0 || feature_errors[i] > point_num_limit) {continue; }
    countfound++;
}
if (countfound <= 0) {countfound = 1;}
PointImg1 = cvCreateMat(countfound, 2, CV_32F);
PointImg2 = cvCreateMat(countfound, 2, CV_32F);
countfound = 0;
matched_features = 0;
for (int i = 0; i < corner_count; i++)
{
    if (features_found[i] == 0 || feature_errors[i] > point_num_limit) {continue; }
    CvPoint p0 = cvPoint(cvRound(cornersA[i].x), cvRound(cornersA[i].y));
    CvPoint p1 = cvPoint(cvRound(cornersB[i].x), cvRound(cornersB[i].y));
    cvmSet(PointImg1, countfound, 0, p0.x);
    cvmSet(PointImg1, countfound, 1, p0.y);
    cvmSet(PointImg2, countfound, 0, p1.x);
    cvmSet(PointImg2, countfound, 1, p1.y);
    countfound++;
}

```

```

        matched_features++;
    }

```

The output of the above code is then used to calculate the homography between the images:

```

CvMat* H = cvCreateMat(3, 3, CV_32FC1);
// get a singular homography if less than 5 points, but use
// 10 to guarantee that you can at least find 10 of them in
// the next image
if (matched_features >= 10) {
    cvFindHomography(PointImg2, PointImg1, H, CV_RANSAC, 1.0, NULL);
}
else{
    cvSetIdentity(H);
}

```

After this code is executed, H contains the homography between the two images.

B Parameters in OpenCV Methods

The various parameters tuned for the OpenCV methods are as follows (written in order as they appear on [2, 3]):

- **cvGoodFeaturesToTrack**: the goal of this function is to detect interest points in the first image

const CvArr* image black and white image you're calculating the homography *from*

CvArr* eigImage temporary image the same size as **image**

CvArr* tempImage another temporary image the same size as **image**

CvPoint2D32f* corners an array of **CvPoint2D32f** where the location of corner points will be stored

int cornerCount variable where the corner count will be stored

double qualityLevel set to 0.01 in the code; denotes the minimum quality parameter for accepting points in the corner count. Setting this higher means there will be more interest points and setting this lower means there will be fewer interest points.

double minDistance set to 5.0 in the code; specifies the minimum distance between interest points. Setting this lower means there is the possibility for interest points to be clustered together more.

const CvArr* mask = NULL set to NULL in the code; denotes the mask used to determine interest points in the image. I experimented with using different masks (i.e. the whole image, not including a 50-pixel border around the edge of

the image) but it didn't yield good results. I did no formal tests of this with the entire pipeline, so there needs to be more thought put into the use of masks. See ?? for more details.

`int blockSize = 3` set to 3 in the code (the default in OpenCV); I did not play with this parameter, so I'm not sure on the effect it has on the detected features.

`int useHarris=0` set to 1 in the code, so the function uses the Harris detector to detect corners in the image.

`double k = 0.04` set to 0.04, the default value; Used by the Harris detector and I did not adjust this parameter.

- `cvFindCornerSubPix`: the goal of this function is to improve on estimates from the Harris detector in determining the location of interest points in the first image

`const CvArr* image` the first image in black and white (the same image the Harris detector was run)

`CvPoint2D32f* corners` the same array of `CvPoint2D32f` that was passed to `cvGoodFeaturesToTrack` that contains the locations of the detected corners. At the end of the method will contain the updated positions of the interest points.

`int count` the same `cornerCount` variable that was returned by the `cvGoodFeaturesToTrack` method for the number of detected corners

`CvSize win` set to `cvSize(10, 10)` in the code; this defines the search space for the coordinate-refining in the method. I did not try different values of this parameter.

`CvSize zero_zone` set to `cvSize(-1, -1)`; a parameter used to calibrate false negatives, but setting it to `(-1, -1)` does not try to exclude singularities. I did not try different values of this parameter.

`CvTermCriteria criteria` set to `cvTermCriteria(CV_TERMCRIT_ITER|CV_TERMCRIT_EPS, 20, 0.03)` as was set for another application and I did not try different values for this parameter; this specifies termination criteria for the algorithm

- `cvCalcOpticalFlowPyrLK`: this function calculates the optical flow using the Lucas-Kanade method with pyramids between two images

`const CvArr* prev` the image at time t in black and white

`const CvArr* curr` the image at time $t + dt$ in black and white

`const CvArr* prevPyr` buffer pyramid for the first image. Initialized to the same size as the `prev` image.

`const CvArr* currPyr` same as `prevPyr` but for the second image

`const CvPoint2D32f* prevFeatures` the returned `corners` variable returned by the `cvFindCornerSubPix` method; these are the interest points for which the flow will be calculated in the second image

`const CvPoint2D32f* currFeatures` an array of `CvPoint2D32f` types of the same size as `prevFeatures`; will contain the positions of the new points at the end of the method

`int count` the same variable as the output `cornerCount` variable from `cvFindCornerSubPix` that contains the number of feature points

`CvSize winSize` set to `cvSize(30, 30)`; defines the size of the window in which to look for the new position of the interest point. Setting this higher will give the ability to potentially find optical flow between two images that are less overlapping, but will also lead to more errors. I tried different values of the window size, but not against any formal tests.

`int level` set to 3; the number of pyramid layers to use. I did not adjust this parameter.

`char* status` an empty array needs to be passed of the size of the maximum number of features possibly found; keeps track of which features were found in the second image

`float* track_error` an empty array needs to be passed the size of the maximum number of features possibly found; keeps track of the error in detection.

`CvTermCriteria criteria` set to `cvTermCriteria(CV_TERMCRIT_ITER | CV_TERMCRIT_EPS, 20, 0.03)`. I did not try different parameters.

`int flags` set to 0.

- `cvFindHomography`: Determines the homography matrix between two images and normalizes it to scale so that the 3,3 entry is 1.

`const CvMat* srcPoints` The concatenated matrix of points output from `prevFeatures` output from `cvCalcOpticalFlowPyrLK`. The code for transforming from `prevFeatures` to `srcPoints` is given in A.

`const CvMat* dstPoints` The same as `srcPoints` except for `currFeatures`.

`CvMat* H` A 3 by 3 matrix that will store the calculated homography.

`int method=0` Set to `CV_RANSAC`; use the RANSAC method to calculate the homography to get rid of outliers in the detected points.

`double ransacReprojThreshold=0` Set to 1.0; used for the RANSAC threshold in the algorithm.

`CvMat* status=NULL` Set to NULL; I did not experiment with this parameter.